



Infinite Horizon MDP

Realisé par:

Mohamed EL YESSEFI

Yahya YOUNES

2A-R

2022-2023

Contents

1	The optimality equation	3
2	Solving the equation by value iteration	4
3	Find the optimal policy by Policy Iteration	5

1 The optimality equation

L'équation d'optimalité : $V^*(s) = \max(R(s, a) + \gamma * \sum P(s'|s, a) * V^*(s'))$

$$V^*(0) = \max(10 + \gamma * V^*(1), 1 + \gamma * V^*(0))$$

$$V^*(1) = \max(0 + \gamma * V^*(1), -15 + \gamma * V^*(0))$$

pour $\gamma = 0$:

$$V^*(0) = 10 + \gamma * V^*(1) \Rightarrow V^*(0) = 10$$

$$V^*(1) = 0 + \gamma * V^*(1) \Rightarrow V^*(1) = 0$$

pour $\gamma = 1$:

$$V^*(0) = 1 + \gamma * V^*(0) \Rightarrow V^*(0) = \frac{1}{1-\gamma}$$

$$V^*(1) = -15 + \gamma * V^*(0) \Rightarrow V^*(1) = -15 + \frac{\gamma}{1-\gamma}$$

Résolution analytique

$$10 = \left(\frac{1}{1-\gamma}\right) \Rightarrow \gamma = 0.9$$

$$0 = -15 + \left(\frac{\gamma}{1-\gamma}\right) \Rightarrow \gamma = 0.9375$$

2 Solving the equation by value iteration

The python code that uses the value iteration :

+ Code + Texte

```
[35] 1 import numpy as np
      2
      3 # This is the function that solves the optimality equation through value iteration.
      4 def iterative_value_evaluation (P, R, discount_factor, epsilon=1e-4):
      5     # Initialize value function V
      6     V = np.zeros (n_states)
      7
      8     # Value iteration
      9     while True:
     10         V_prev = V. copy ()
     11         for s in range(n_states):
     12             V[s] = max([R[s,a] + discount_factor * np. sum(P[s, :] * V_prev) for a in range(n_states)])
     13         if np.abs(V - V_prev).max() < epsilon :
     14             break
     15
     16     # Calculate optimal policy
     17     policy = np.zeros (n_states, dtype=int)
     18
     19     for s in range(n_states):
     20         policy[s] = np. argmax( [R[s, a] + discount_factor * np.sum(P[s, :] * V) for a in range(n_states)])
     21
     22     return(V,policy)

[37] 1 # Number of states
      2 n_states = 2
      3
      4 # Number of actions
      5 n_actions = 2
      6
      7 # Transition matrix
      8 P = np.array ( [[0, 1], [0.9, 0.1]])
      9
     10 # Reward matrix
     11 R = np. array ( [[10, 1], [-15, 0]])
     12 # Test the function
     13 for discount_factor in [0, 0.9, 0.999]:
     14     V, policy = iterative_value_evaluation (P, R, discount_factor, epsilon=1e-4)
     15     print(f"Discount Factor = {discount_factor}: Optimal Policy = {policy}, Values = {V}")
     16
```

Figure 1: A screenshot of our code on python

The results we get :

```

➤ Discount Factor = 0: Optimal Policy = [0 1], Values = [10.  0.]
Discount Factor = 0.9: Optimal Policy = [0 1], Values = [50.27541761 44.75055573]
Discount Factor = 0.999: Optimal Policy = [0 1], Values = [4739.51366834 4734.24801619]

```

Figure 2: A screenshot of the result we get when running the python code

3 Find the optimal policy by Policy Iteration

We implemented this Python code that uses iteration policy to find the optimal policy for different values of γ

```

✓ [77] 1 import numpy as np
0s      2
        3 # Number of states
        4 n_states = 2
        5
        6 # Number of actions
        7 n_actions = 2
        8
        9 # Transition matrix
       10 P = np.array([[10, 1], [0.9, 0.1]])
       11
       12 # Reward matrix
       13 R = np.array([[10, 1], [-15, 0]])
       14
       15 # Function that evaluates a policy using value iteration.
       16 def iterative_policy_evaluation(P, R, gamma, epsilon=1e-4):
       17     # Initialization of V
       18     V = np.zeros(n_states)
       19
       20     # Value iteration
       21     while True:
       22         V_prev = V.copy()
       23         for s in range(n_states):
       24             V[s] = sum([policy[s] * (R[s, a] + gamma * np.dot(P[:, s, a], V_prev)) for a in range(n_actions)])
       25             if np.abs(V - V_prev).max() < epsilon:
       26                 break
       27     return V
       28
       29 def iterative_policy_improvement(P, R, V, gamma):
       30     # Function that improves a policy using policy iteration
       31     # Initialization of the policy
       32     policy = np.zeros(n_states, dtype=int)
       33
       34     # Policy improvement
       35     for s in range(n_states):
       36         policy[s] = np.argmax([R[s, a] + gamma * np.dot(P[:, s, a], V) for a in range(n_actions)])
       37     return policy
       38

```

```
38
39
40 def iterative_policy_iteration(P, R, gamma, epsilon=1e-4):
41     # Function that improves a policy using policy iteration
42     # Initialization of the policy
43     policy = np.zeros(n_states, dtype=int)
44
45     # Policy iteration
46     while True:
47         # Policy evaluation
48         V = iterative_policy_evaluation(P, R, gamma, epsilon)
49
50         # Policy improvement
51         policy_prev = policy.copy()
52         policy = iterative_policy_improvement(P, R, V, gamma)
53
54         # Stop if the policy converges
55         if (policy == policy_prev).all():
56             break
57     return V, policy
58
```

Figure 3: A screenshot of the python code